



PUPPET DESIGN PATTERNS

*David Danzilio
Kovarus, Inc.*



ABOUT ME

- Architect with Kovarus
- Joined in May 2016
- Previously at Constant Contact, Sandia National Laboratories, and a few other places you've never heard of
- Operations background, but more of a developer these days
- One of the maintainers of Vox Pupuli
- Organizer of the Boston Puppet User Group



danzilio/virtualbox

A module to install VirtualBox on Linux hosts.

APPROVED

Version 1.7.0 • Apr 7, 2016 • 5,171 downloads

31,590 | 5.0



danzilio/letsencrypt

A module to install the Letsencrypt client and request certs.

Version 1.0.0 • Feb 22, 2016 • 5,737 downloads

8,060 | 5.0



danzilio/kickstart

A module to generate Kickstart configurations with Puppet.

Version 0.4.0 • Apr 6, 2016 • 847 downloads

3,810 | 5.0



danzilio/report_all_the_things

Dumps the complete Puppet report for processing by a report processor.

Version 0.1.0 • Jul 13, 2015 • 3,136 downloads

3,136 | 4.5



danzilio/scribe_reporter

A report processor that sends whole Puppet reports to a Scribe log server.

Version 0.2.0 • Jul 14, 2015 • 2,431 downloads

2,579 | 4.5

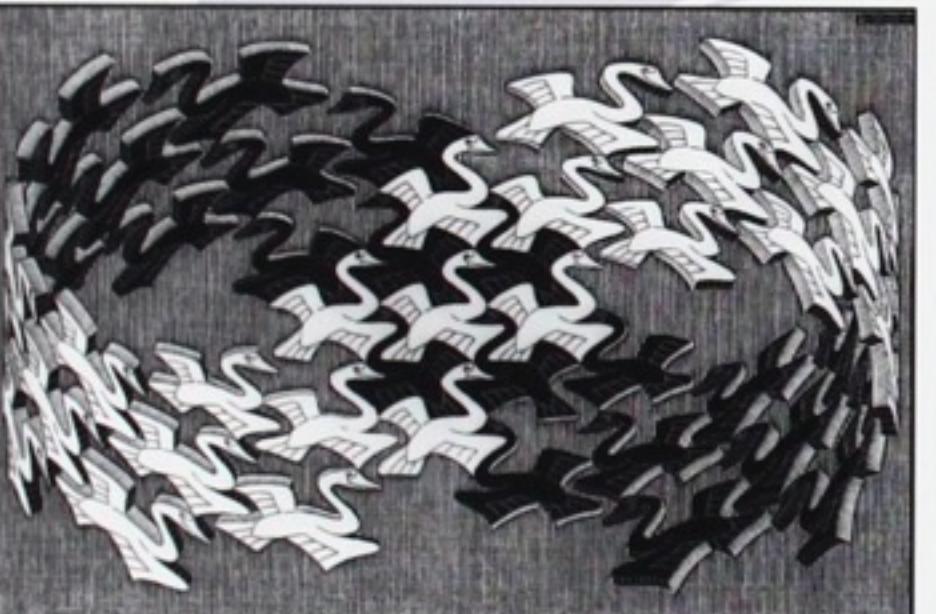


ABOUT DESIGN PATTERNS

Design Patterns

Elements of Reusable
Object-Oriented Software

Erich Gamma
Richard Helm
Ralph Johnson
John Vlissides



Cover art © 1994 M.C. Escher / Cordon Art - Baarn - Holland. All rights reserved.

Foreword by Grady Booch

ADDISON-WESLEY PROFESSIONAL COMPUTING SERIES



GANG OF FOUR

- Gang of Four (GoF) book introduced the concept for Object Oriented programming
- 23 patterns with examples written in C++ and SmallTalk
- Published in 1994, more than 500,000 copies sold
- One of the best selling software engineering books in history
- Influenced an entire generation of developers, languages, and tools

DESIGN PATTERNS

- Can be highly contextual and language dependent, but a lot can be learned from all of them
- The GoF focused on statically-typed, object oriented, compiled languages
- Some languages have implemented primitives for these patterns
- Not many of the GoF patterns directly apply to Puppet
- All of the GoF patterns focus on reinforcing a set of design principles

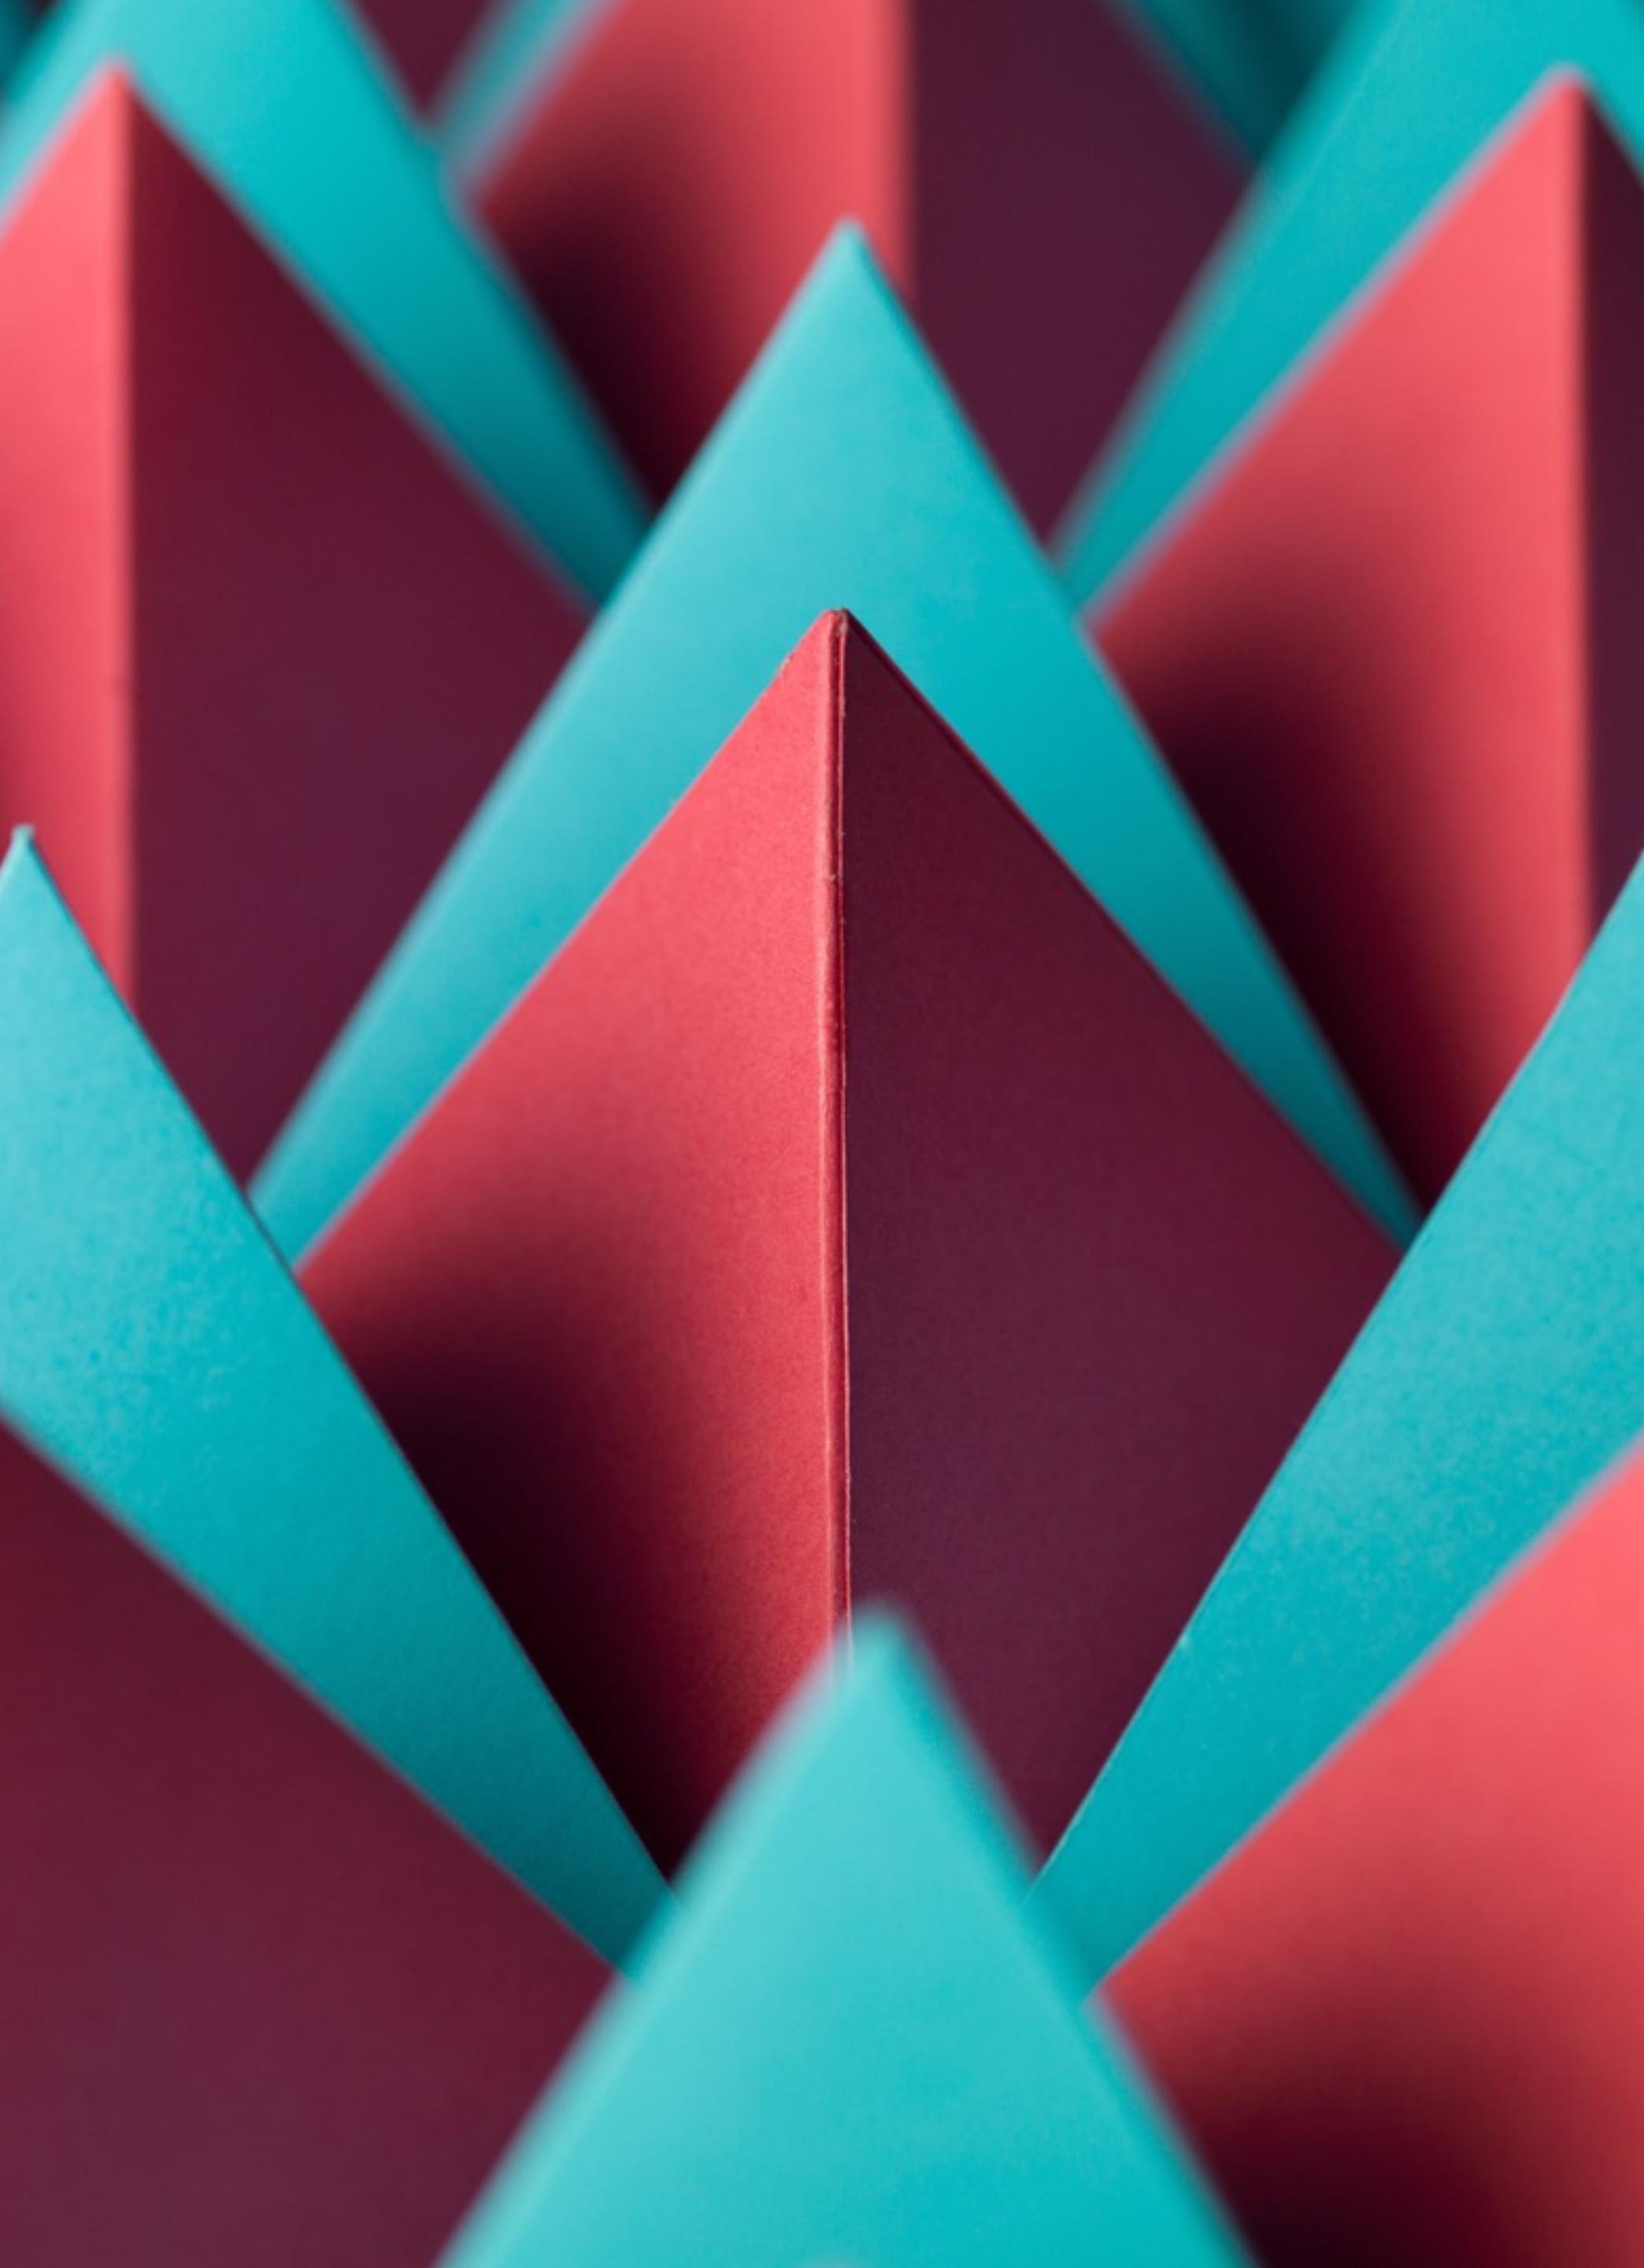
WHY DESIGN PATTERNS?

DESIGN PRINCIPLES

SEPARATE THINGS THAT CHANGE

Minimize risk when making changes





LOOSE COUPLING

Reduce dependencies, increase modularity

SEPARATION OF CONCERNS

Divide your application into distinct features



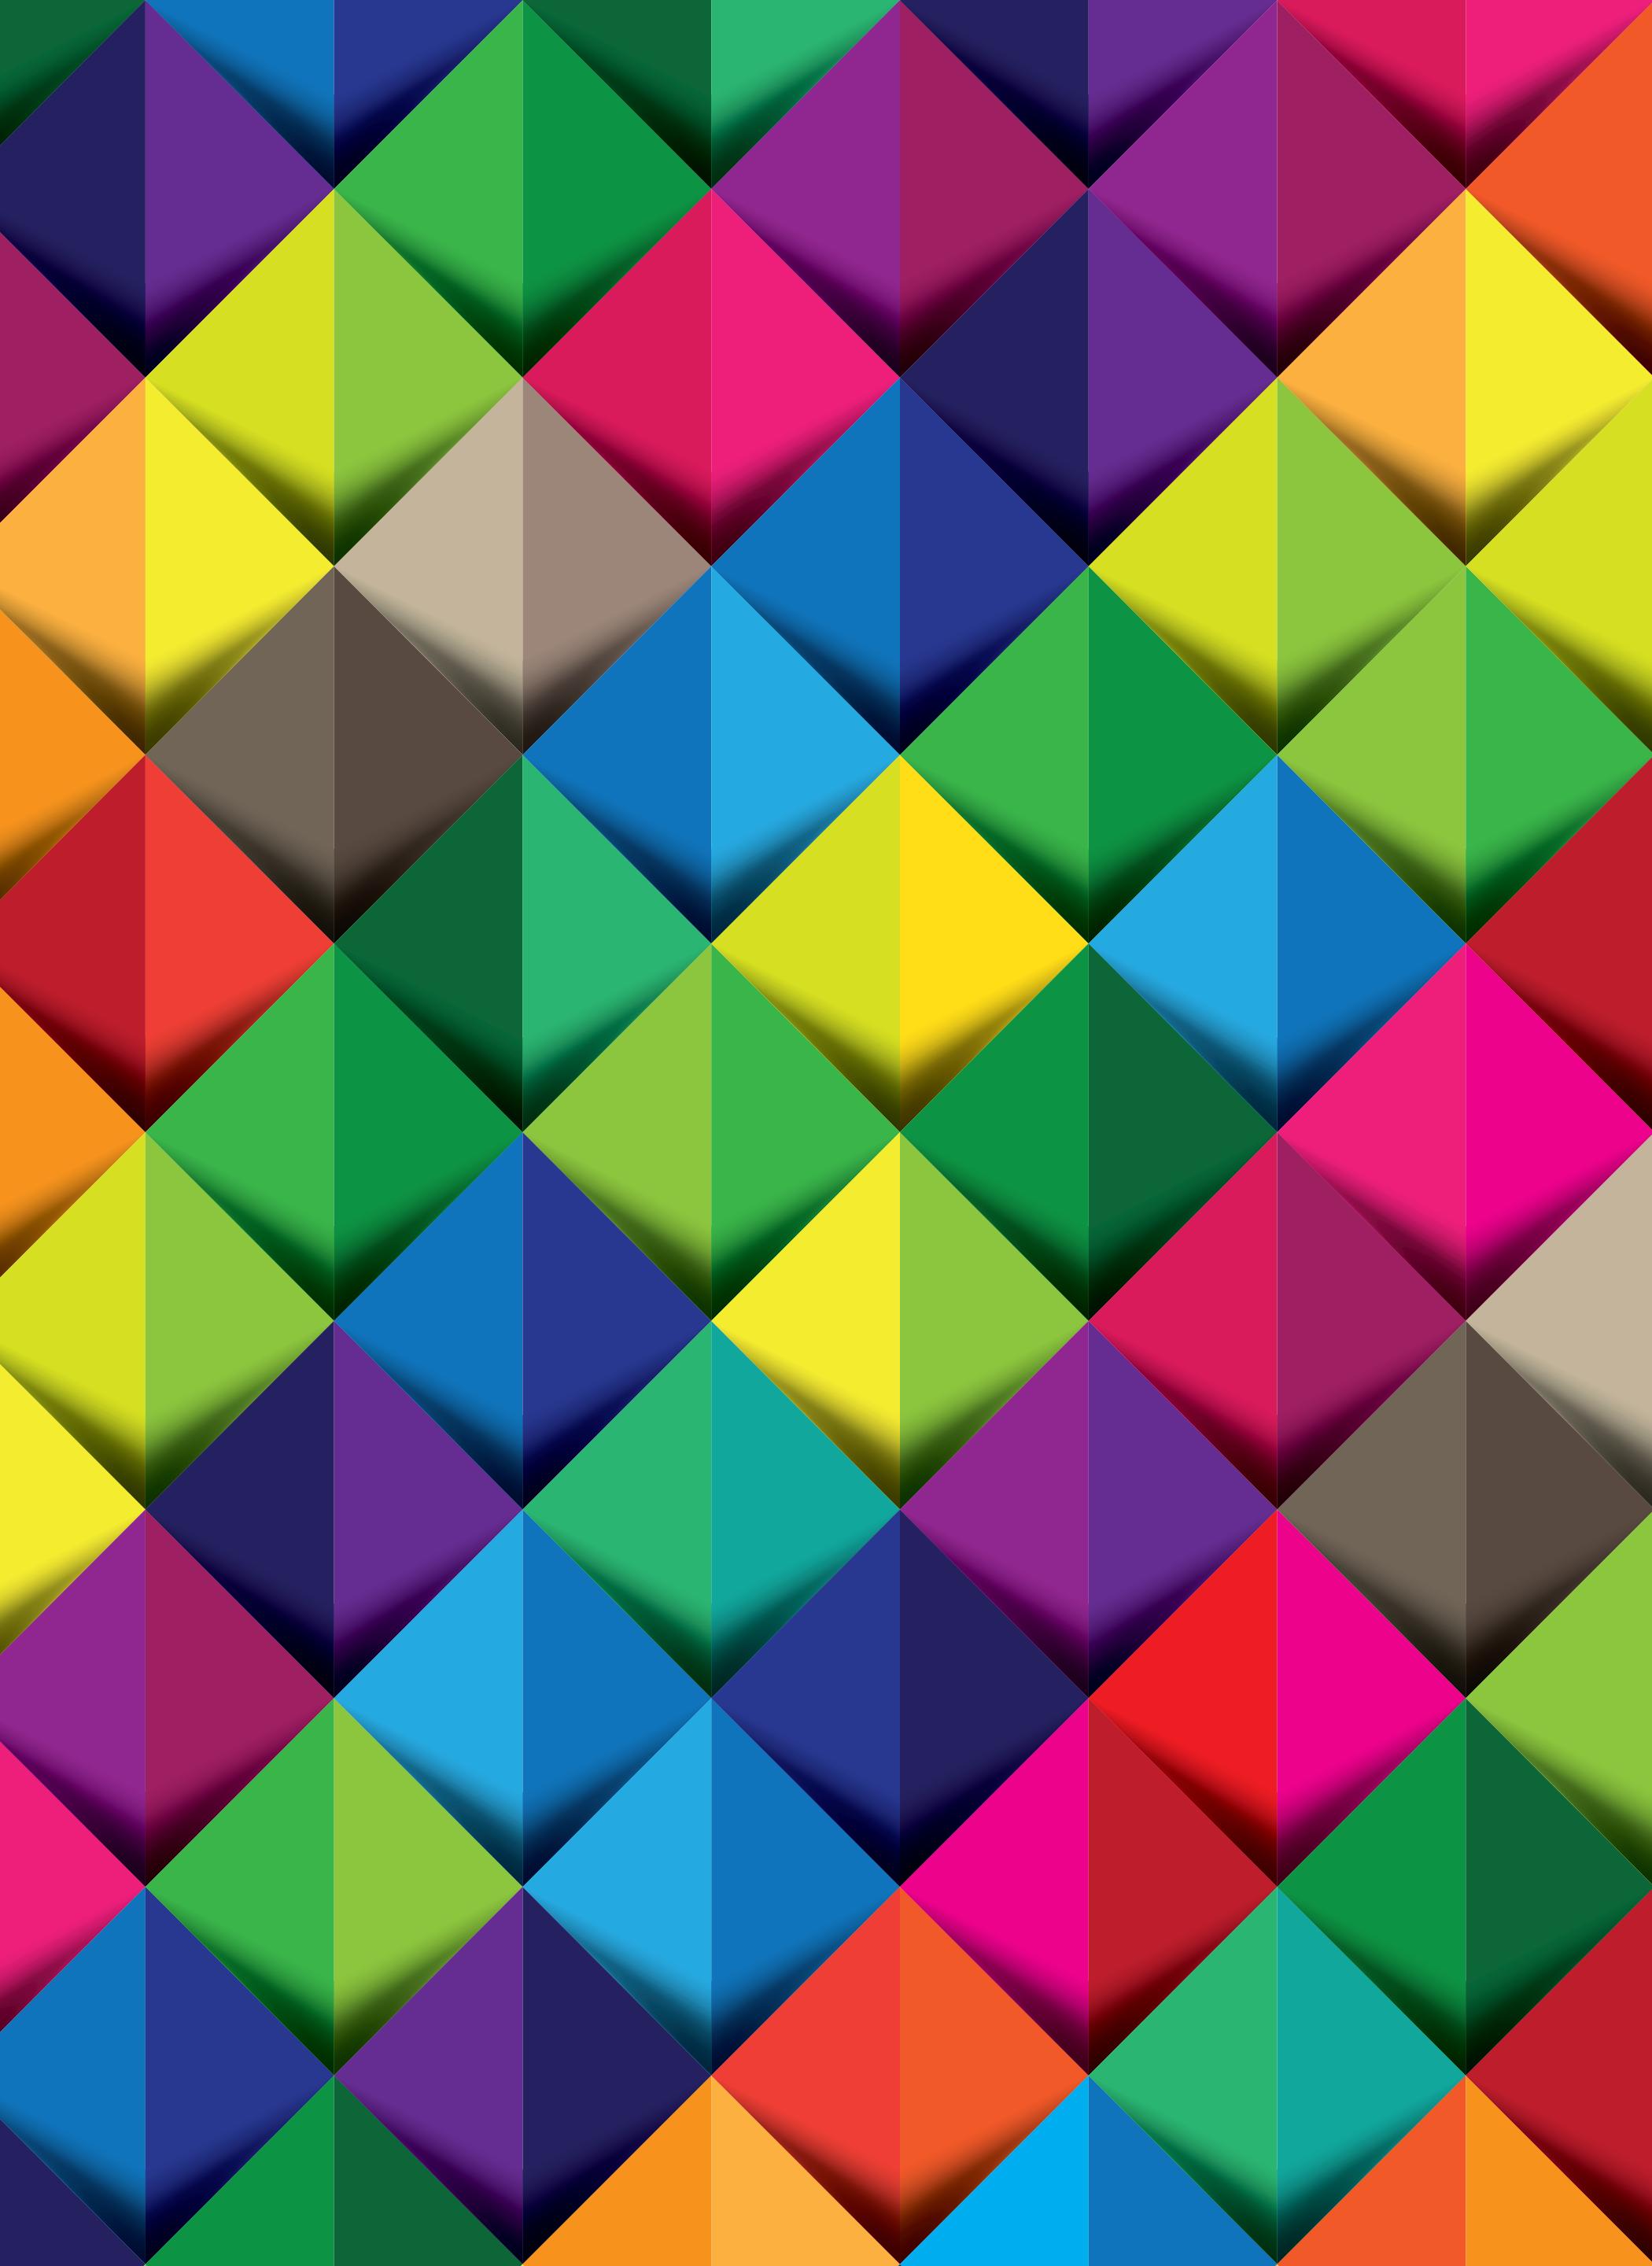


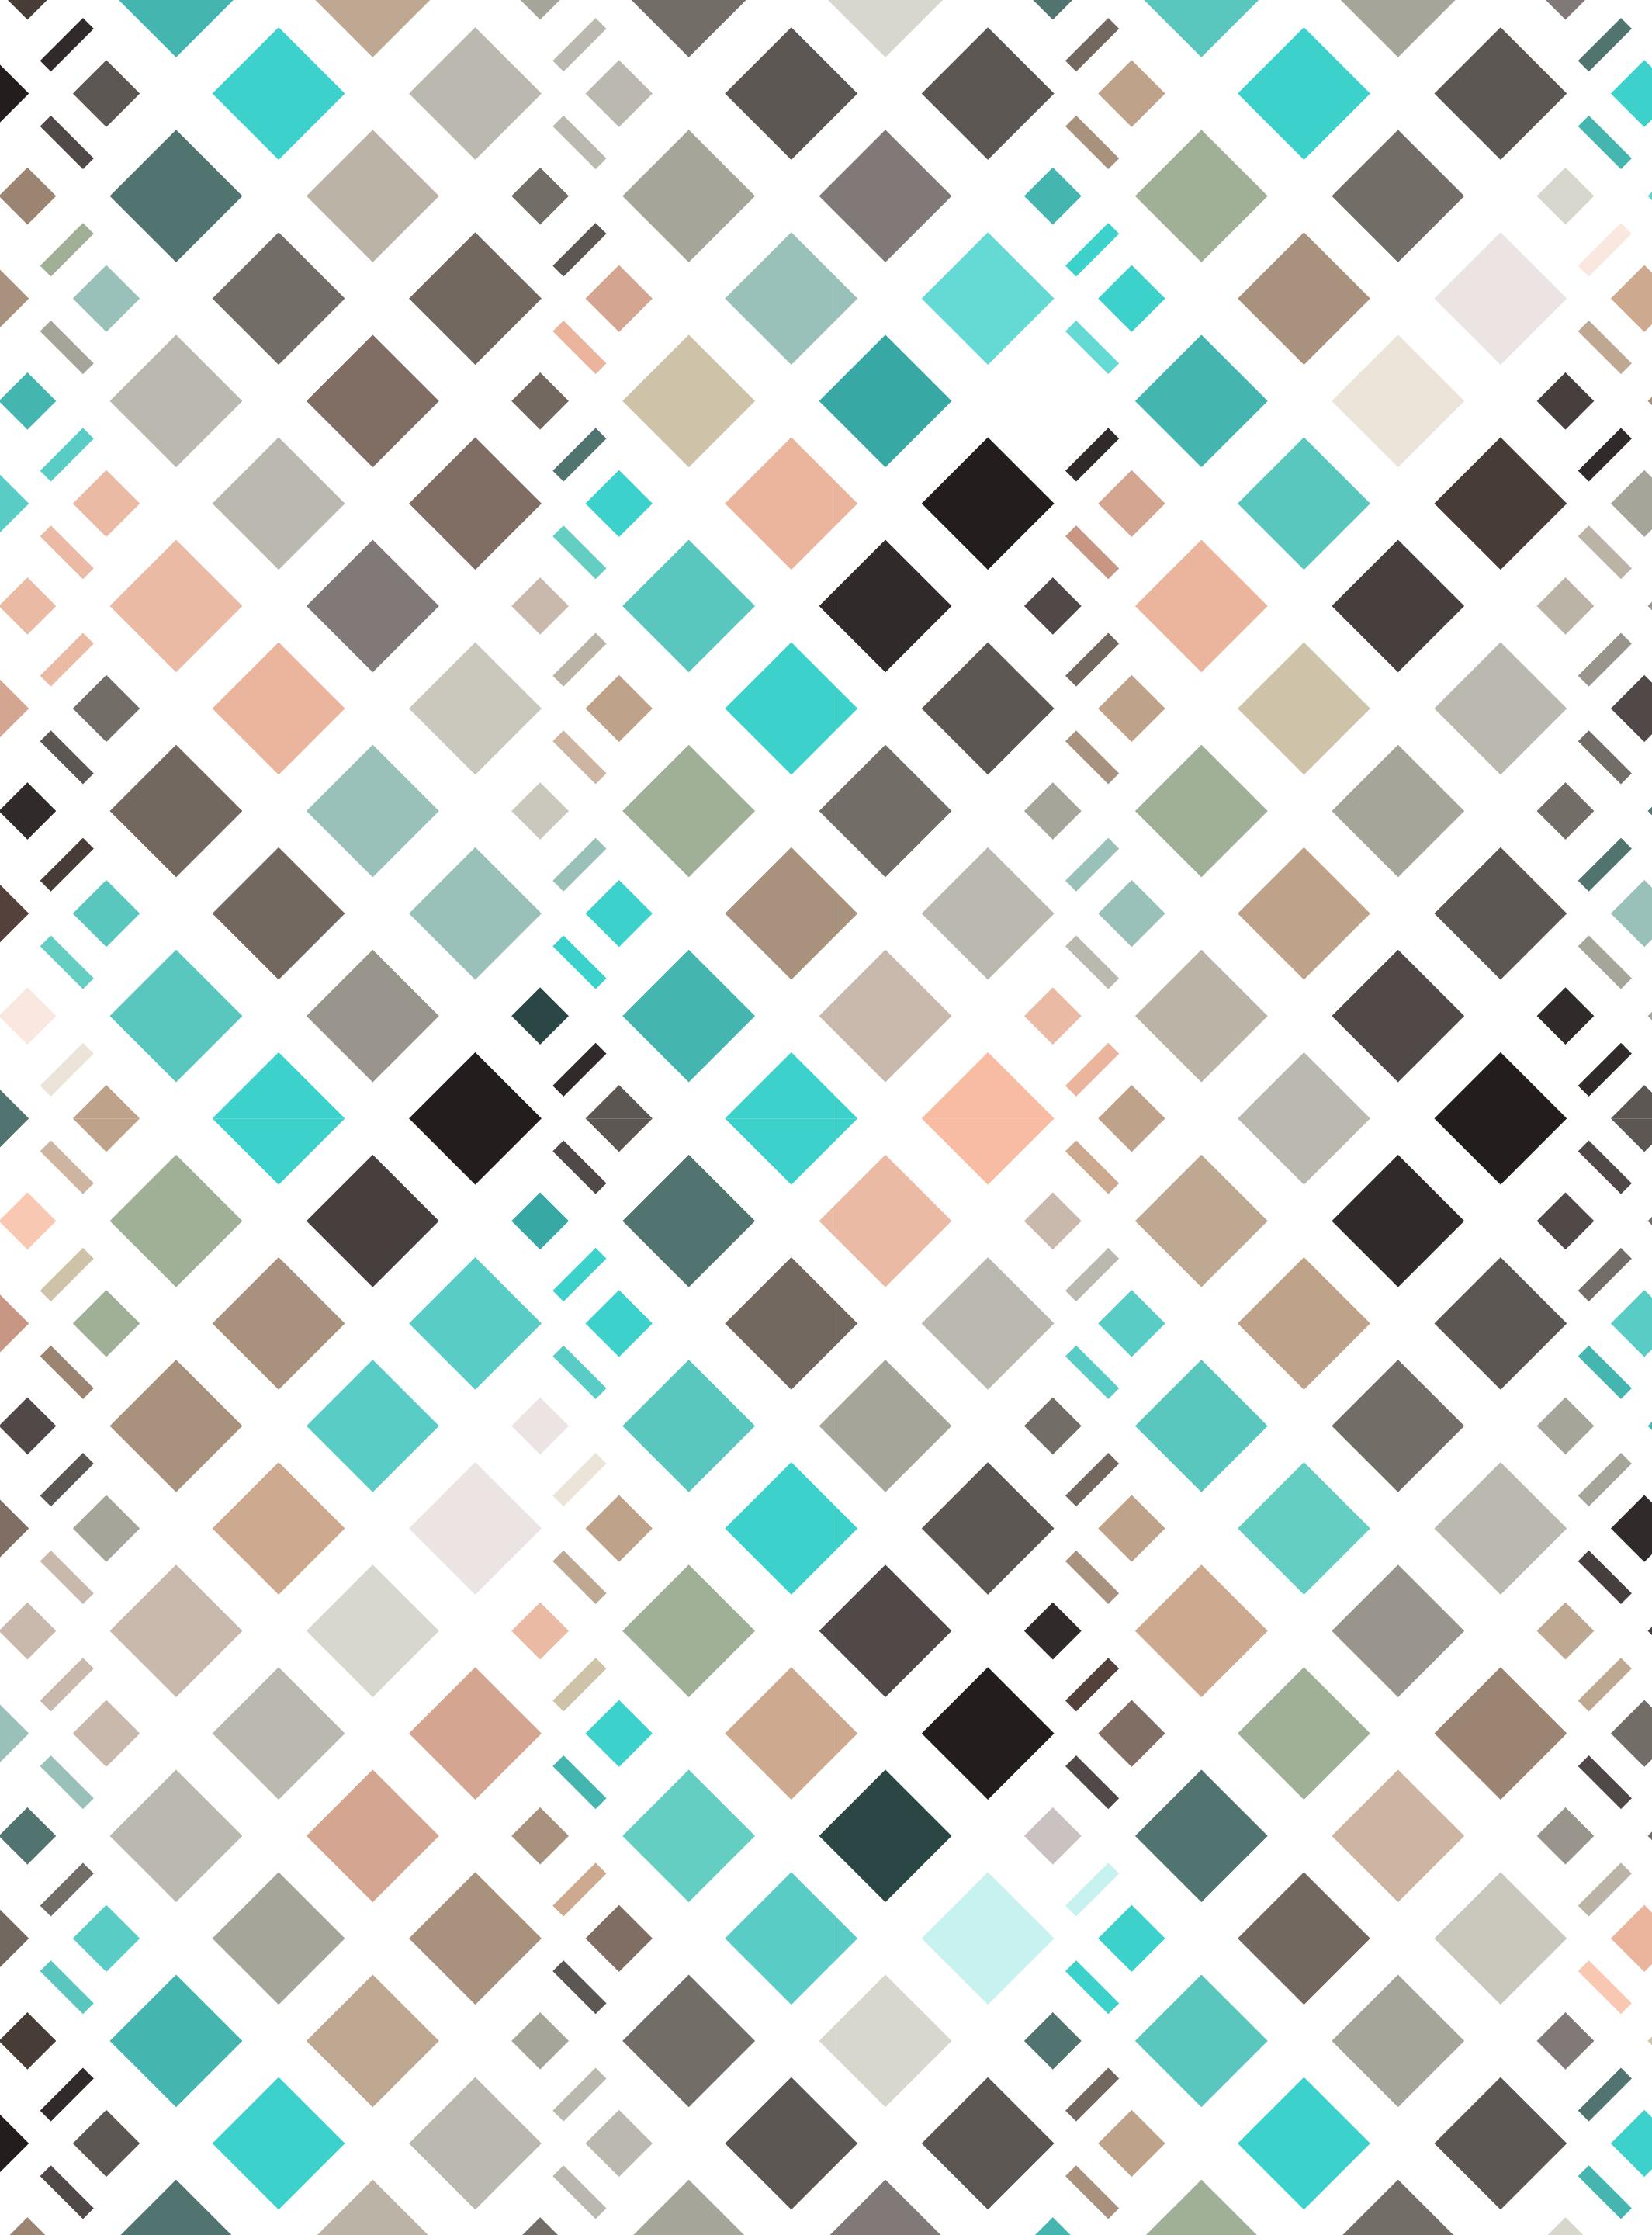
SINGLE RESPONSIBILITY

Do one thing well

LEAST KNOWLEDGE

Components should know as little as possible about their neighbors



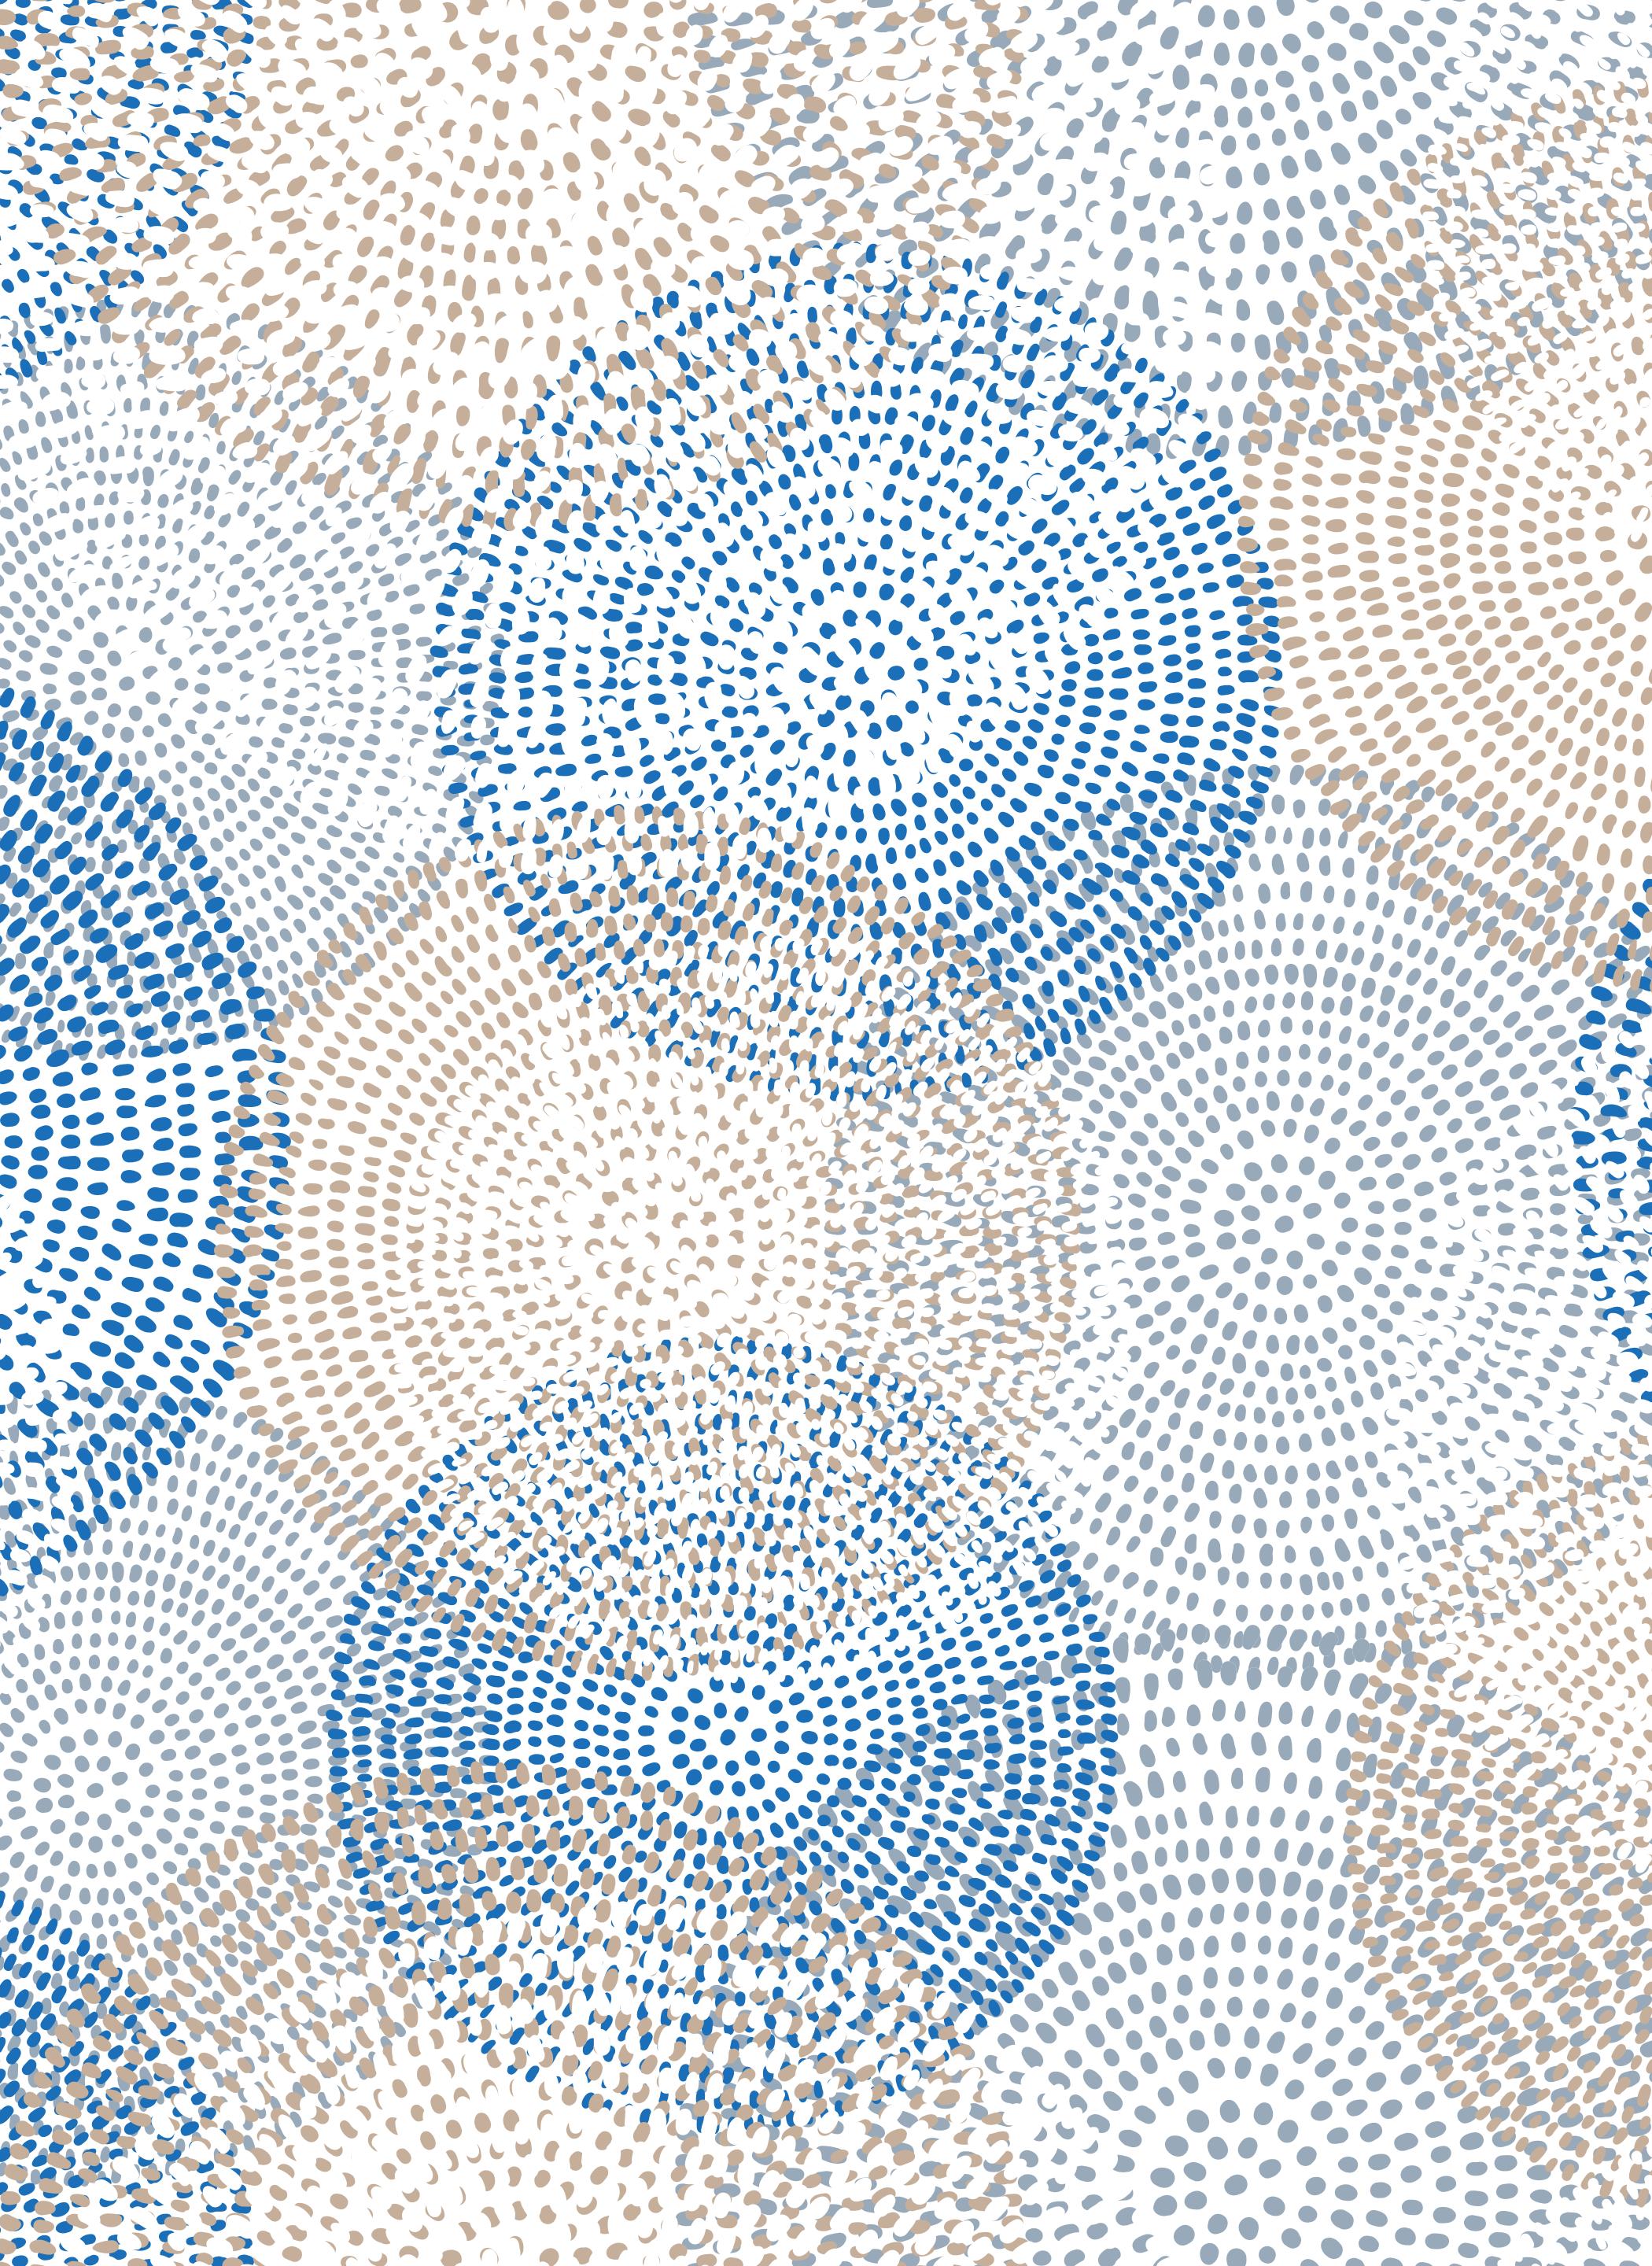


DON'T REPEAT YOURSELF

DRY, because we're lazy

PROGRAM TO AN INTERFACE

*Interfaces are more stable than the
implementation*





SIX PATTERNS

SIX PATTERNS

- **Resource Wrapper:** adding functionality to code you don't own
- **Package, File, Service:** breaking up monolithic classes
- **Params:** delegating parameter defaults
- **Strategy:** doing the same thing differently
- **Roles and Profiles:** organizing your code for modularity
- **Factory:** creating resources on the fly



THE RESOURCE WRAPPER PATTERN

Extending existing resources

ABOUT THE WRAPPER PATTERN

- **Problem:** resources you don't own are missing some functionality or feature required to implement your requirements
- **Solution:** use composition to add your required functionality without modifying the code you don't own

ABOUT THE WRAPPER PATTERN

- Use the resource wrapper pattern when you need to **add functionality to an existing resource**
- When you feel the need to write your own resources, or to make changes to Forge modules, think about whether **you should really be using the wrapper pattern**
- You can do this in Puppet 3 and Puppet 4, but **it's much cleaner in Puppet 4**
- This pattern forms the basis of many other patterns you'll see today

EXAMPLE: MANAGING USERS

- We want to manage our employee user accounts
- Requirements:
 - The user's UID should be set to their employee ID
 - All employees need to be members of the 'employees' group
 - We should manage a user's bash profile by default, but users may opt out of this upon request

EXAMPLE: MANAGING USERS

Employee ID is used for the user ID →

```
define mycompany::user (
  $employee_id,
  $gid,
  $groups      = ['employees'],
  $username     = $title,
  $manage_profile = true,
) {
```

All employees should be in the 'employees' group →

```
if $manage_profile {
  file { "/home/${username}/.bash_profile":
    ensure => file,
    owner  => $username,
    require => User[$username],
  }
}
```

Feature flag to manage your user's bash profile →

```
user { $username:
  uid    => $employee_id,
  gid    => $gid,
  groups => $groups,
}
```

Manage ~/.bash_profile with a file resource →

Pass your parameters to the user resource →

EXAMPLE: MANAGING USERS

“We have a new employee named Bob. He’s employee 1093 and he needs to be a member of the wheel group so he can sudo. He wants to manage his own bash profile.”

```
mycompany::user { 'bob':  
  employee_id    => '1093',  
  gid            => 'wheel',  
  manage_profile => false,  
}
```

EXAMPLE: MANAGING USERS

“Hey, I’d like to have my shell set to /bin/zsh, can you do that for me?”

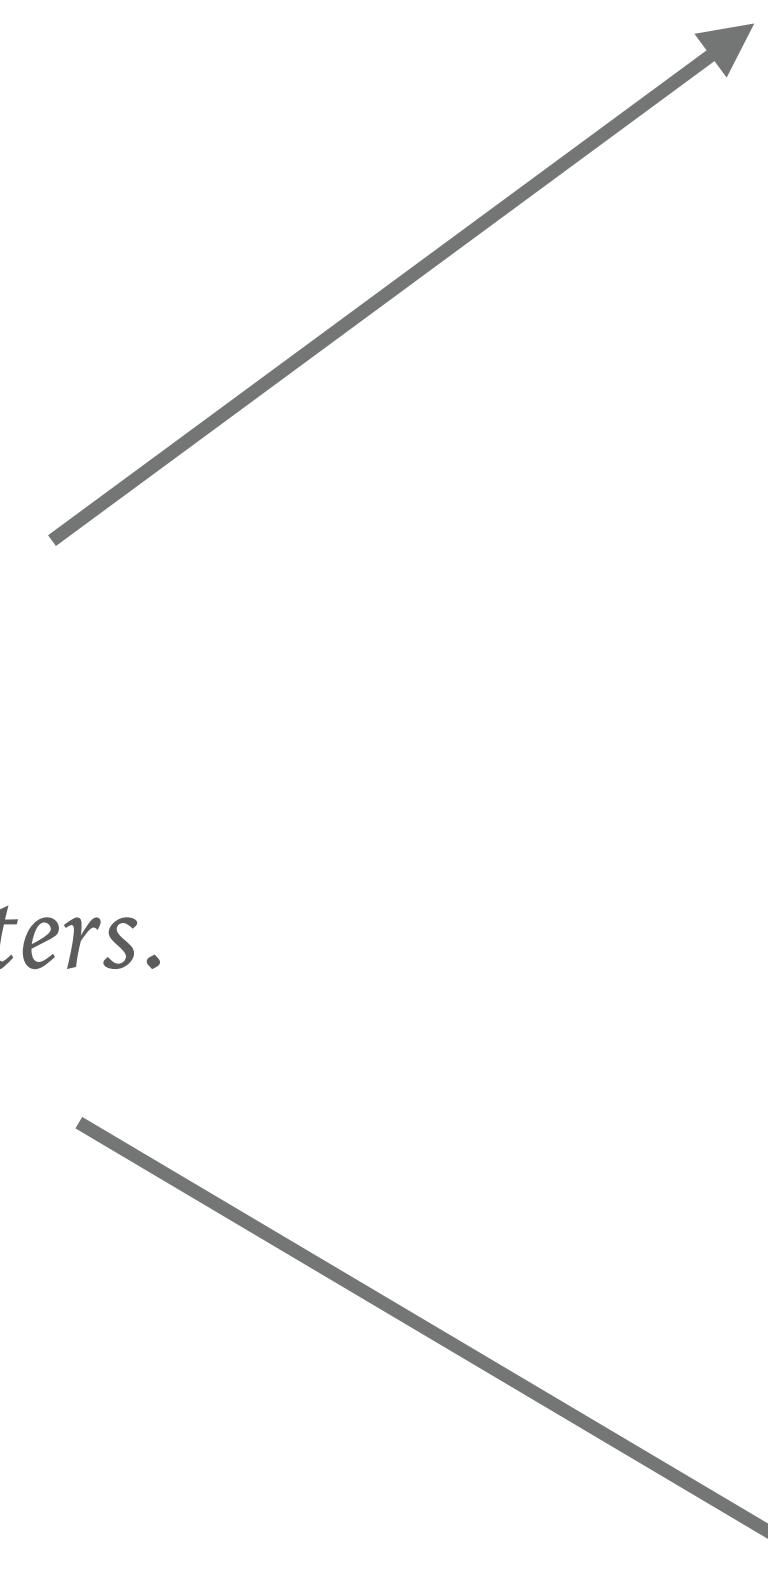
```
mycompany::user { 'bob':  
  employee_id    => '1093',  
  gid            => 'wheel',  
  shell          => '/bin/zsh',  
  manage_profile => false,  
}
```

*Could not retrieve catalog:
Invalid parameter 'shell' for
type 'mycompany::user'*

EXAMPLE: MANAGING USERS

Problem

You must maintain these parameters.



```
define mycompany::user (
  $employee_id,
  $gid,
  $groups      = ['employees'],
  $username    = $title,
  $manage_profile = true,
) {

  if $manage_profile {
    file { "/home/${username}/.bash_profile":
      ensure => file,
      owner  => $username,
      require => User[$username],
    }
  }

  user { $username:
    uid    => $employee_id,
    gid    => $gid,
    groups => $groups,
  }
}
```

EXAMPLE: MANAGING USERS (PUPPET 3)

Much more flexible interface

```
define mycompany::user (
  $username      = $title,
  $manage_profile = true,
  $user          = {}
) {

  if $manage_profile {
    file { "/home/${username}/.bash_profile":
      ensure  => file,
      owner   => $username,
      require => User[$username],
    }
  }
}
```

Enforce business rules

```
$user_defaults = { 'groups' => ['employees'] }
$user_params = merge($user, $user_defaults)
```

*Create the user resource by passing
the hash to create_resources*

```
create_resources('user', $username, $user_params)
}
```

EXAMPLE: MANAGING USERS (PUPPET 4)

Much more flexible interface

```
define mycompany::user (
  String  $username      = $title,
  Boolean $manage_profile = true,
  Hash    $user          = {}
)

if $manage_profile {
  file { "/home/${username}/.bash_profile":
    ensure  => file,
    owner   => $username,
    require => User[$username],
  }
}

$user_defaults = { 'groups' => ['employees'] }

user { $username:
  * => $user_defaults + $user,
}
```

Enforce business rules

*Use splat operator to pass hash keys
as parameters to the user resource*

EXAMPLE: MANAGING USERS

“Hey, I’d like to have my shell set to /bin/zsh, can you do that for me?”

```
mycompany::user { 'bob':  
  employee_id    => '1093',  
  gid            => 'wheel',  
  manage_profile => false,  
}  
  
→
```

```
mycompany::user { 'bob':  
  manage_profile => false,  
  user           => {  
    'uid'          => '1093',  
    'gid'          => 'wheel',  
  }  
}
```

EXAMPLE: MANAGING USERS

“Hey, I’d like to have my shell set to /bin/zsh, can you do that for me?”

```
mycompany::user { 'bob':  
  employee_id    => '1093',  
  gid            => 'wheel',  
  manage_profile => false,  
}  
  
→
```

```
mycompany::user { 'bob':  
  manage_profile => false,  
  user           => {  
    'uid'          => '1093',  
    'gid'          => 'wheel',  
    'shell'        => '/bin/zsh',  
  }  
}
```

THE PACKAGE/ FILE/SERVICE PATTERN

Assembling complex behavior



ABOUT THE PACKAGE/FILE/SERVICE PATTERN

- **Problem:** your module's `init.pp` is getting too cluttered because all of your code lives in that one file
- **Solution:** break out the basic functions of your module into separate classes, generally into a package, config, and service class

ABOUT THE PACKAGE/FILE/SERVICE PATTERN

- This is one of the first patterns you see when learning Puppet
- This is the embodiment of the **Single Responsibility** and **Separation of Concerns** principles
- Most modules can be broken down into some form of Package, Config File, and Service management
- Use this specific pattern **any time you write a module** that manages these things
- Keep the spirit of this pattern in mind **whenever you write a module** that is more than a few lines long
- This has the added benefit of allowing us to utilize **class containment** for cleaner resource ordering

EXAMPLE: NTP

Set some variables based on the osfamily fact

```
class ntp {  
  case $::osfamily {  
    'Solaris': {  
      $package_name = ['SUNWntpr', 'SUNWntpu']  
      $config_file = '/etc/inet/ntp.conf'  
      $service_name = 'network/ntp'  
    }  
    'RedHat': {  
      $package_name = 'ntp'  
      $config_file = '/etc/ntp.conf'  
      $service_name = 'ntpd'  
    }  
  }  
}
```

Installs the ntp package for that platform

```
  package { $package_name:  
    ensure => installed,  
  }
```

Places the ntp config file

```
  file { $config_file:  
    ensure => file,  
    content => template('ntp/ntp.conf.erb'),  
    require => Package[$package_name],  
    notify => Service[$service_name],  
  }
```

*Ensures that the package is installed first
Notifies the ntp service of changes to the file*

Manages the ntp service

```
  service { $service_name:  
    ensure => running  
  }
```

EXAMPLE: NTP

```
class ntp {  
    case $osfamily {  
        'Solaris': {  
            $package_name = ['SUNWntpr', 'SUNWntpu']  
            $config_file = '/etc/inet/ntp.conf'  
            $service_name = 'network/ntp'  
        }  
        'RedHat': {  
            $package_name = 'ntp'  
            $config_file = '/etc/ntp.conf'  
            $service_name = 'ntpd'  
        }  
    }  
}
```

Installs the ntp package → class { 'ntp::install': } -> ← *Package is installed first*
Places the ntp config file → class { 'ntp::config': } ~> ← *Notifies the service of changes*
Manages the ntp service → class { 'ntp::service': } ←

EXAMPLE: NTP

```
class ntp::install {
  package { $ntp::package_name:
    ensure => installed,
  }
}

class ntp::config {
  file { $ntp::config_file:
    ensure => file,
    content => template('ntp/ntp.conf.erb'),
  }
}

class ntp::service {
  service { $ntp::service_name:
    ensure => running,
  }
}
```



THE PARAMS PATTERN

Separating out your data

ABOUT THE PARAMS PATTERN

- **Problem:** hard-coded data makes modules fragile, verbose parameter default and variable setting logic make classes hard to read
- **Solution:** convert embedded data to parameters and move that data to a separate class where it can be used as parameter defaults

ABOUT THE PARAMS PATTERN

- The params pattern **breaks out your variable assignment and parameter defaults** into a separate class, typically named params
- Makes classes **easier to read** by moving the default setting logic into a purpose-built class
- **Delegates responsibility** for setting defaults to the params class
- Module data is a new feature designed to eliminate the params pattern by **moving this logic into Hiera**
- Until module data becomes more ubiquitous, you'll see params in use in **almost every module**
- Use this pattern **any time you have data that must live in your module**

EXAMPLE: NTP

Problems

*Only supports Solaris and RedHat
~70% of this class is devoted to data*

```
class ntp {  
    case $osfamily {  
        'Solaris': {  
            $package_name = ['SUNWntpr', 'SUNWntpu']  
            $config_file = '/etc/inet/ntp.conf'  
            $service_name = 'network/ntp'  
        }  
        'RedHat': {  
            $package_name = 'ntp'  
            $config_file = '/etc/ntp.conf'  
            $service_name = 'ntpd'  
        }  
    }  
  
    class { 'ntp::install': } ->  
    class { 'ntp::config': } ~>  
    class { 'ntp::service': }  
}
```

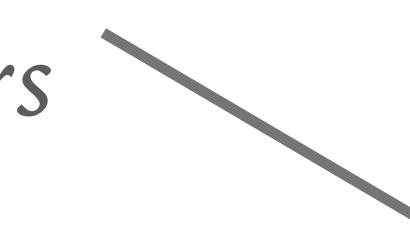
EXAMPLE: NTP

*Create a purpose-built class
to store your module's data*



```
class ntp::params {  
  case $::osfamily {  
    'Solaris': {  
      $package_name = ['SUNWntpr', 'SUNWntp']  
      $config_file = '/etc/inet/ntp.conf'  
      $service_name = 'network/ntp'  
    }  
    'RedHat': {  
      $package_name = 'ntp'  
      $config_file = '/etc/ntp.conf'  
      $service_name = 'ntpd'  
    }  
  }  
}
```

*These variables become the default
values for your class parameters*



EXAMPLE: NTP

Convert the variables to parameters, and set their defaults to the corresponding variables in the params class

Inheriting the params class ensures that it is evaluated first

```
class ntp (
    $package_name = $ntp::params::package_name,
    $config_file = $ntp::params::config_file,
    $service_name = $ntp::params::service_name,
) inherits ntp::params {

    class { 'ntp::install': } ->
    class { 'ntp::config': } ~>
    class { 'ntp::service': }
}
```

THE STRATEGY PATTERN

Varying the algorithm



ABOUT THE STRATEGY PATTERN

- **Problem:** you have multiple ways to achieve basically the same thing in your module, but you need to choose one way based on some criteria (usually a fact)
- **Solution:** break each approach into separate classes and let your caller decide which to include

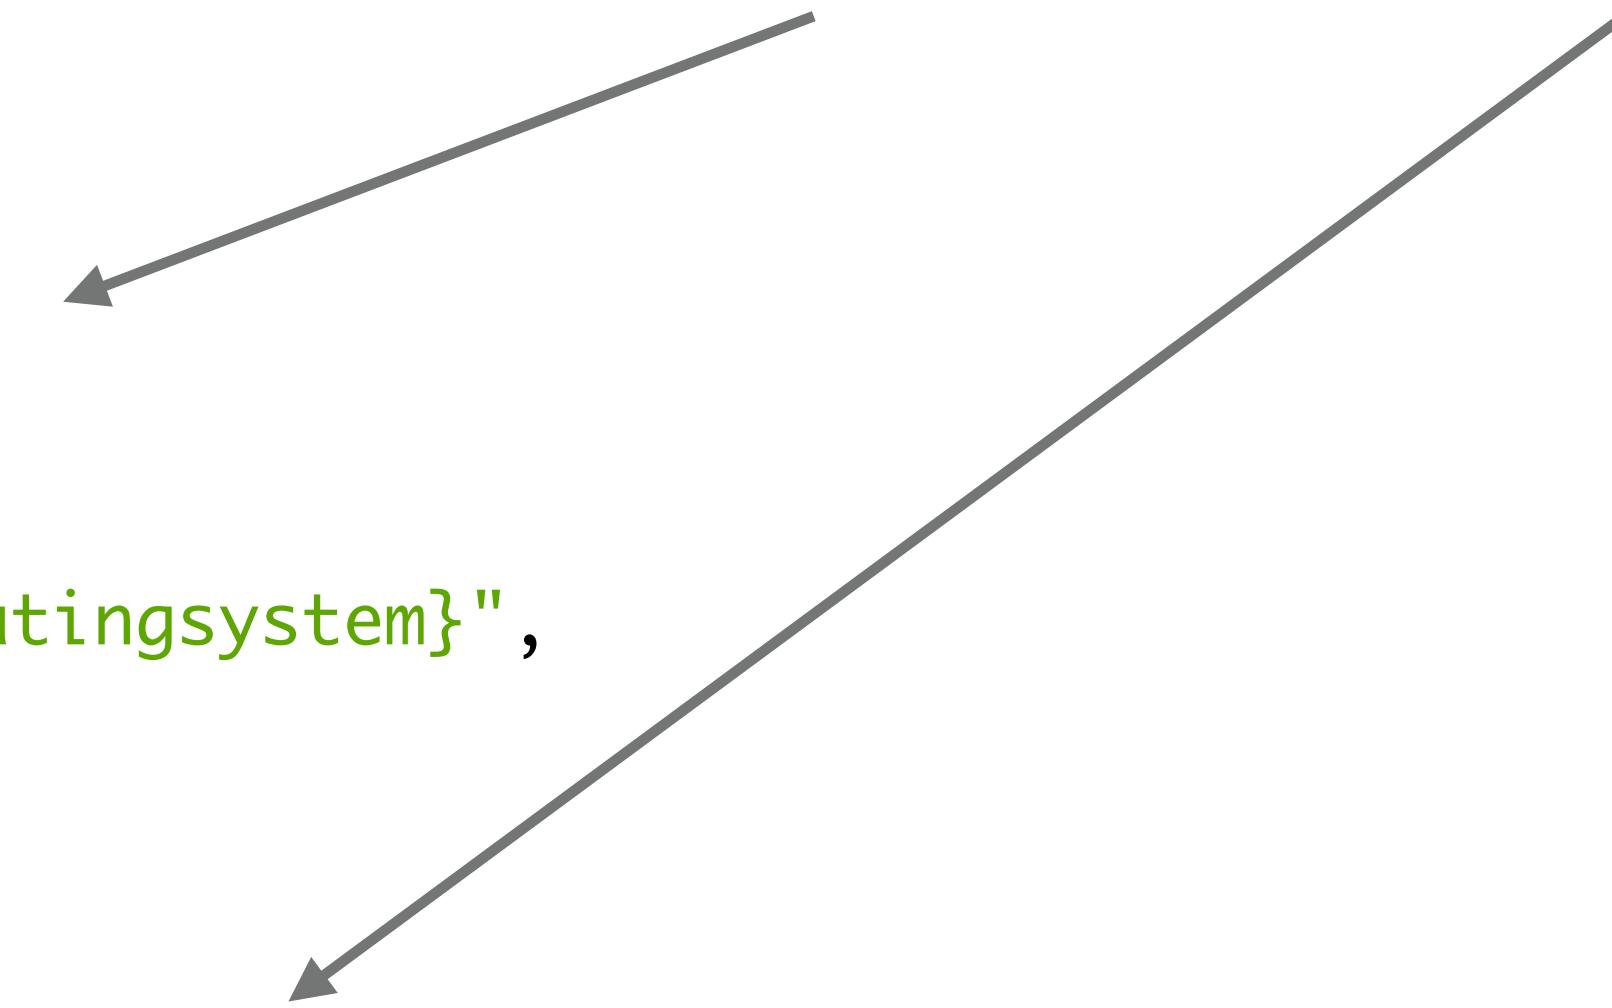
ABOUT THE STRATEGY PATTERN

- This is a GoF pattern
- Use this pattern when you have lots of logic doing effectively the same thing but with different details under certain conditions
- The Strategy Pattern uses **composition** to assemble complex behavior from smaller classes

EXAMPLE: MYSQL

```
class mysql {  
    ...  
  
    case $::osfamily {  
        'Debian': {  
            apt::source { 'mysql':  
                comment  => 'MySQL Community APT repository',  
                location => "http://repo.mysql.com/apt/$::operatingsystem",  
                release   => $::lsbdistcodename,  
                repos     => 'mysql-5.7',  
                include   => { src => false },  
            }  
        }  
        'RedHat': {  
            yumrepo { 'mysql':  
                descr   => 'MySQL Community YUM repository',  
                baseurl => "http://repo.mysql.com/yum/mysql-5.7-community/el/$::lsbmajdistrelease/$::architecture",  
                enabled  => true,  
            }  
        }  
    }  
    ...  
}
```

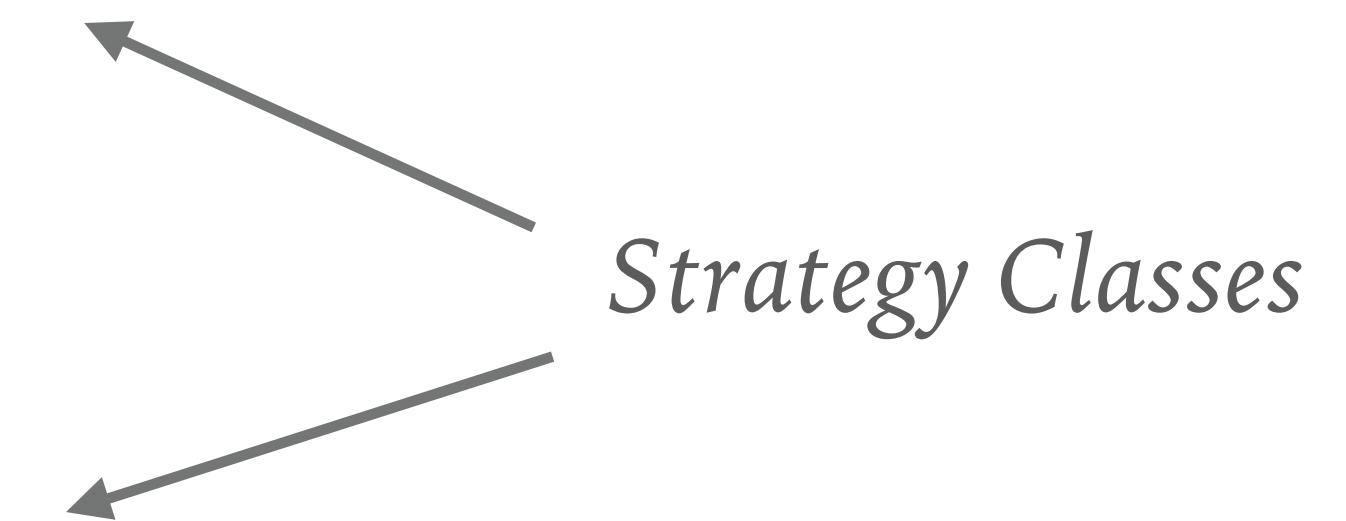
Both managing a package repository



EXAMPLE: MYSQL

```
class mysql::repo::debian {
  apt::source { 'mysql':
    comment => 'MySQL Community APT repository',
    location => "http://repo.mysql.com/apt/${::operatingsystem}",
    release => $::lsbdistcodename,
    repos => 'mysql-5.7',
    include => { src => false },
  }
}
```

```
class mysql::repo::redhat {
  yumrepo { 'mysql':
    descr => 'MySQL Community YUM repository',
    baseurl => "http://repo.mysql.com/yum/mysql-5.7-community/el/${::lsbmajdistrelease}/${::architecture}",
    enabled => true,
  }
}
```



Strategy Classes

EXAMPLE: MYSQL

Context Class → `class mysql {`

Case statement determines which strategy class to include → `case $::osfamily {`

```
class mysql {  
    ...  
  
    case $::osfamily {  
        'Debian': { include mysql::repo::debian }  
        'RedHat': { include mysql::repo::redhat }  
    }  
    ...  
}
```



THE ROLES AND PROFILES PATTERN

Organizing your code for modularity

ABOUT THE ROLES AND PROFILES PATTERN

- **Problem:** large node statements with many classes, lots of inherited node statements, difficulty identifying what a server's purpose is in your environment
- **Solution:** add an extra layer of abstraction between your node and your modules

ABOUT THE ROLES AND PROFILES PATTERN

- The Roles and Profiles pattern was described by **Craig Dunn** in his blog post *Designing Puppet - Roles and Profiles*
- This is one of the **most comprehensive design patterns** for Puppet
- It is the “official” way to structure your Puppet code
- You should **always use Roles and Profiles**
- Craig does an excellent job describing these concepts in depth, you should read his blog post here: <http://www.craigdunn.org/2012/05/239/>

WITHOUT ROLES AND PROFILES

Base node includes common modules

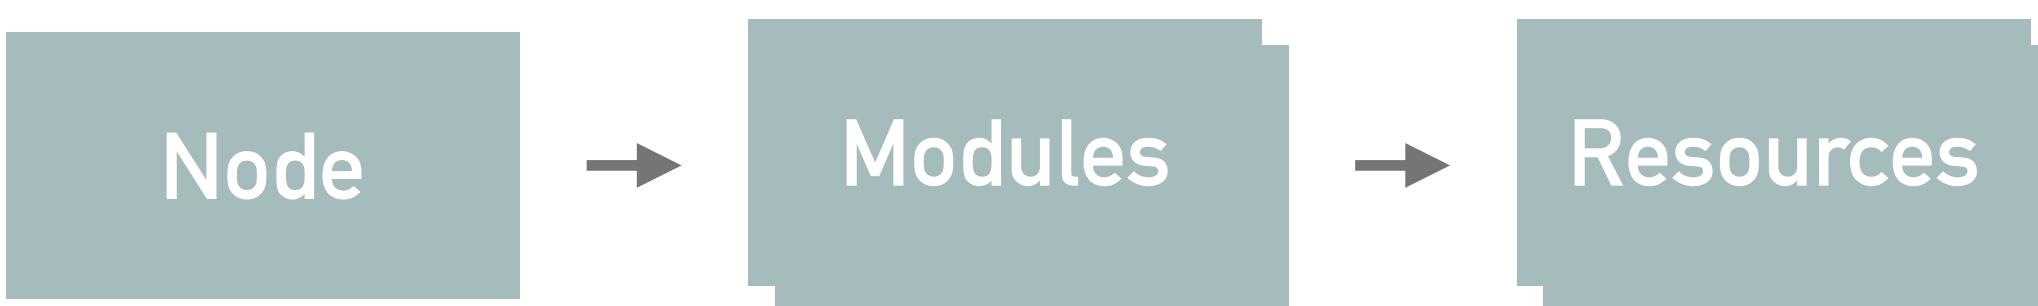
Nodes inherit the base to get common functionality

Problems

This file can get really long, really fast

Can violate DRY when you have lots of similar nodes

Edge cases are hard to manage



More specific functionality is added in each node statement

```
node base {  
  include mycompany::settings  
}
```

```
node www inherits base {  
  include apache  
  include mysql  
  include php  
  include nagios::web_server  
}
```

```
node ns1 inherits base {  
  include bind  
  include nagios::bind  
}
```

```
bind::zone { 'example.com': type => master }  
}
```

```
node ns2 inherits base {  
  include bind  
  include nagios::bind  
}
```

```
bind::zone { 'example.com': type => slave }  
}
```

ROLES AND PROFILES TERMINOLOGY

- Module: implements **one piece of software** or functionality
- Profile: combines modules to **implement a stack** (i.e. “A LAMP stack includes the apache, mysql, and php modules”)
- Role: combine profiles to **implement your business rules** (i.e. “This server is a web server”)



- A node can only ever include **one role**
- If you think you need to include two roles, you've probably just identified another role

CONVERTING TO ROLES AND PROFILES

```
class roles::base {
  include profiles::base
}

class roles::web_server {
  include profiles::base
  include profiles::lamp
}

class roles::nameserver::master inherits roles::base {
  include profiles::bind::master
}

class roles::nameserver::slave inherits roles::base {
  include profiles::bind::slave
}
```

```
class profiles::base {
  include mycompany::settings
}

class profiles::lamp {
  include apache
  include mysql
  include php
  include nagios::web_server
}

class profiles::bind ($type = master) {
  include bind
  bind::zone { 'example.com':
    type => $type,
  }
}

class profiles::bind::master {
  include profiles::bind
}

class profiles::bind::slave {
  class { 'profiles::bind':
    type => slave,
  }
}
```

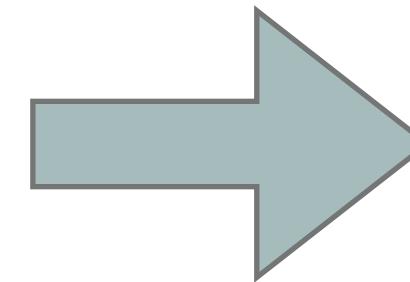
CONVERTING TO ROLES AND PROFILES

```
node base {  
  include mycompany::settings  
}
```

```
node www inherits base {  
  include apache  
  include mysql  
  include php  
  include nagios::web_server  
}
```

```
node ns1 inherits base {  
  include bind  
  include nagios::bind  
  bind::zone { 'example.com': type => master }  
}
```

```
node ns2 inherits base {  
  include bind  
  include nagios::bind  
  bind::zone { 'example.com': type => slave }  
}
```



```
node www {  
  include roles::web_server  
}  
  
node ns1 {  
  include roles::nameserver::master  
}  
  
node ns2 {  
  include roles::nameserver::slave  
}
```

THE FACTORY PATTERN

Creating resources in your classes



ABOUT THE FACTORY PATTERN

- **Problem:** your module has to create a lot of resources of the same type, or you want to control how resources are created with your module
- **Solution:** create the resources in your class based on data passed to your parameters

ABOUT THE FACTORY PATTERN

- This is also known as the `create_resources` pattern
- Emerged early on as crude iteration support in older Puppet versions
- We already saw this in action in the **Resource Wrapper Pattern** example
- Use this pattern when you want your module to have a **single entry point**, even for creating your own resources

EXAMPLE: MANAGING CONFIGURATION WITH INI_SETTING

```
class puppet (
  $ca_server  = 'puppet-ca.example.com',
  $master     = 'puppet.example.com',
  $pluginsync = true,
  $noop       = false,
) {  
  
  $defaults = { 'path' => '/etc/puppet/puppet.conf' }  
  $main_section = {  
    'main/ca_server' => { 'setting' => 'ca_server', 'value' => $ca_server },  
    'main/server'     => { 'setting' => 'server', 'value' => $master },  
  }  
  
  $agent_section = {  
    'agent/pluginsync' => { 'setting' => 'pluginsync', 'value' => $pluginsync },  
    'agent/noop'        => { 'setting' => 'noop', 'value' => $noop },  
  }  
  
  create_resources('ini_setting', $main_section, merge($defaults, { section => 'main' }))  
  create_resources('ini_setting', $agent_section, merge($defaults, { section => 'agent' }))  
}
```

Get data from params

Organize the data so we can consume it with create_resources

Pass the munged data to create_resources

EXAMPLE: MANAGING CONFIGURATION WITH INI_SETTING (PUPPET 4)

```
class puppet {
  String $path          = '/etc/puppet/puppet.conf',
  Hash  $main_section  = {
    'ca_server'  => 'puppet-ca.example.com',
    'server'      => 'puppet.example.com'
  },
  Hash  $agent_section = {
    'pluginsync' => true,
    'noop'        => false,
  },
}

['agent', 'main'].each |$section| {
  $data = getvar("${section}_section")
  $data.each |$key,$val| {
    ini_setting { "${section}/${key}":
      path    => $path,
      section => $section,
      setting => $key,
      value   => $val,
    }
  }
}
```

Pass a hash for each section

Iterate over each section name

Fetch the variable that holds that section's data

Iterate over that data, passing it to an ini_setting resource

THE END

STAY TUNED FOR MORE PATTERNS

CONTACT INFO



- [@djdanzilio](https://twitter.com/djdanzilio) on Twitter
- [danzilio](#) on Freenode (you can usually find me in #voxpupuli, #puppet-dev, and #puppet)
- [danzilio](#) on GitHub and The Forge
- [ddanzilio \(at\) kovarus \(dot\) com](mailto:ddanzilio@kovarus.com)
- blog.danzil.io
- www.kovarus.com